

# The Reader-Writer Problem

## Implementation and Analysis of Synchronization Strategies

Your Name

Your University

January 13, 2026

# Outline

1 Introduction

2 System Architecture

3 Implementation

# The Reader-Writer Problem

## Problem Statement

Multiple concurrent threads accessing shared resources:

- **Readers**: Only read data (can operate concurrently)
- **Writers**: Modify data (require exclusive access)

## Challenge

Allow multiple readers **OR** one writer at a time, while maintaining:

- **Correctness**: No data corruption
- **Performance**: Maximize throughput
- **Fairness**: Prevent starvation

# Real-World Applications

## Database Systems

- PostgreSQL
- MySQL InnoDB
- Read/Write transactions

## Programming Languages

- Java ReadWriteLock
- C++ shared\_mutex
- Rust RwLock

## Operating Systems

- Linux kernel RW semaphores
- File system access
- Process synchronization


## Distributed Systems

- Cache coherence
- Shared memory
- Configuration services

# Research Contributions

- ① **Unified Framework:** Single API supporting 4 synchronization modes
- ② **Comprehensive Testing:** 16 automated runs demonstrating race conditions and correctness
- ③ **Quantitative Analysis:** Performance trade-offs between strategies
- ④ **Practical Insights:** Implementation guidance for real systems

# Three-Layer Architecture



## Benefits:

- Easy mode comparison (single parameter change)
- Reusable synchronization primitives
- Configurable test scenarios

# Unified Lock API

```
typedef enum {
    VANILLA,           // No synchronization
    READER_PREF,       // Reader preference
    WRITER_PREF,       // Writer preference
    FAIR               // Fair scheduling
} rw_mode_t;

// Unified API
void rw_init(rw_lock_t *lock, rw_mode_t mode);
void reader_enter(rw_lock_t *lock);
void reader_exit(rw_lock_t *lock);
void writer_enter(rw_lock_t *lock);
void writer_exit(rw_lock_t *lock);
void rw_destroy(rw_lock_t *lock);
```

**Key advantage:** Performance differences purely due to synchronization strategy, not implementation variations

# Synchronization Primitives

## Mutex-Based Implementation (Our Choice)

- Uses `pthread_mutex_t`
- Better integration with condition variables
- Clearer ownership semantics
- Widely supported in POSIX

## Alternative: Semaphore-Based

- Uses counting semaphores
- Simpler conceptual model for some algorithms
- Common in textbook solutions

**Note:** Both approaches are equally valid. Reader/Writer Preference and Fair scheduling are **algorithms**, not primitives.

# Shared String Application

## Shared Resource

char shared\_string[256] - A mutable string buffer

## Writer Behavior:

- ① Select sentence from 20 predefined strings
- ② Acquire writer lock
- ③ Copy sentence character-by-character
- ④ Release lock
- ⑤ Log operation

## Reader Behavior:

- ① Acquire reader lock
- ② Read entire string
- ③ Release lock
- ④ Validate against valid set
- ⑤ Log operation

## Race Condition: Torn Reads

Without synchronization, readers see partial updates:

"Syncating systems manage..." (mixed sentences)

# Mode 1: Vanilla (No Synchronization)

```
void reader_enter(rw_lock_t *lock) {  
    // No synchronization!  
    lock->active_readers++;  
}
```

## Purpose

Baseline to demonstrate race conditions

## Expected Behavior

- Data corruption (torn reads)
- Lost updates
- **DO NOT USE IN PRODUCTION!**

## Mode 2: Reader Preference

### Algorithm

- First reader locks resource
- Subsequent readers increment counter (no blocking)
- Last reader unlocks resource
- Writers wait for all readers to finish

### Advantages:

- Maximizes read throughput
- Multiple concurrent readers
- Low reader latency

### Disadvantages:

- **Writer starvation**
- Continuous readers block writers indefinitely

**Use case:** Read-heavy workloads with infrequent writes

# Mode 3: Writer Preference

## Algorithm

- Writers acquire `read_try` lock
- Blocks new readers when writers waiting
- Existing readers finish, then writer executes
- Readers wait for all writers to complete

## Advantages:

- Prevents writer starvation
- Ensures timely updates
- Data freshness

## Disadvantages:

- Reader starvation
- Continuous writers delay readers

**Use case:** Write-heavy workloads requiring fresh data

# Mode 4: Fair Scheduling (Turnstile Pattern)

## Algorithm

- All threads pass through queue\_lock "turnstile"
- FIFO ordering - no cutting in line
- Both readers and writers get fair access
- Prevents starvation of either type

## Advantages:

- No starvation
- Balanced access
- Predictable latency

## Disadvantages:

- Slight throughput reduction
- Extra lock overhead

**Use case:** Mixed workloads requiring fairness guarantees

# Comparison Summary

Mode	Correctness	Starvation	Throughput
Vanilla	NO	N/A	High
Reader Pref	YES	Writers	Highest
Writer Pref	YES	Readers	Medium
Fair	YES	None	Medium

**Key Insight:** No single "best" solution - choice depends on:

- Workload characteristics (read/write ratio)
- Latency requirements
- Fairness constraints

# Test Methodology

## Automated Testing Framework

- **Total runs:** 16 (4 modes × 4 runs)
- **Configuration per run:**
  - 8 writer threads
  - 5 reader threads
  - 8 seconds duration
- **Validation:** 20 predefined valid sentences
- **Detection:** Torn read = any string not in valid set

## Tools

- `run_tests.sh`: Execute tests, save timestamped logs
- `analyze_comprehensive.py`: Parse logs, detect errors, generate report

# Actual Test Results

Session ID: 20260113\_004235

Mode	Clean Runs	Avg Torn Reads
vanilla	0/4	367
reader_pref	4/4	0
writer_pref	4/4	0
fair	4/4	0

## Key Findings

- **Perfect validation:** All synchronized modes 100% correct (12/12 runs)
- **Clear problem demonstration:** Vanilla 0% success rate
- **No false positives:** Zero errors in synchronized runs

# Torn Read Examples (Vanilla Mode)

## Example Corrupted Strings

- ① "Syncating systems manage..."
  - Mixed: "Sync" from one sentence + "ating systems" from another
- ② "Mutal exclusures..."
  - Partial overwrite: "Mutual exclusion" → "Mutal exclusures"
- ③ "Pthread libuduce and..."
  - Character-level corruption from race

**Average:** 367 torn reads per 8-second vanilla run

**Conclusion:** Concurrent writes without synchronization = severe data corruption

# Statistical Significance

## Test Coverage

- **16 total runs** across all modes
- **Hundreds of operations** per run (reads + writes)
- **Multiple sessions** with consistent results

## Success Rates

- Vanilla: 0/4 clean = **0% success** (expected)
- Synchronized modes: 12/12 clean = **100% success**
- High confidence in implementation correctness

**Reproducibility:** Results consistent across multiple test sessions

# Key Findings

## ① Vanilla mode demonstrates the problem

- 100% race condition rate
- Average 367 torn reads per run
- Clear evidence of why synchronization is needed

## ② All synchronized modes achieve correctness

- 100% success rate (12/12 runs)
- Zero data corruption
- Validates implementation

## ③ Trade-offs between strategies

- Reader preference: High throughput, writer starvation
- Writer preference: Data freshness, reader delays
- Fair: Balanced, slight overhead

# Practical Recommendations

## When to Use Each Mode

### Reader Preference:

- Read-heavy workloads (90%+ reads)
- Infrequent writes acceptable to be delayed
- Example: Configuration cache

### Writer Preference:

- Write-heavy workloads
- Fresh data critical
- Example: Real-time monitoring

### Fair Scheduling:

- Mixed workloads
- Fairness guarantees required
- Example: Shared services with SLAs

# Future Directions

- **Advanced mechanisms:**

- Priority-based scheduling
- Adaptive algorithms based on workload

- **Lock-free alternatives:**

- Atomic operations
- Read-Copy-Update (RCU)
- Compare-and-swap patterns

- **Performance optimization:**

- NUMA architecture considerations
- Cache-line alignment
- Fine-grained locking

# Conclusion

## Summary

- Implemented and analyzed 4 Reader-Writer synchronization strategies
- Demonstrated race conditions and validated correctness through comprehensive testing
- Quantified trade-offs between performance, fairness, and starvation

## Contributions

- Unified framework for easy comparison
- Practical implementation insights
- Open-source educational material

**Thank You!**

**Questions?**

